


Standard prompt search recovers no triggers — our scans do

same 10 backdoored models (2 trigger families × 5 scales, refusal)



False alarms on clean models: σ_1 Hessian scan 0/6 · token look-up 1/6 · GCG / RD-GCG never claim a trigger.